

Natural Language Processing - Problem Set #3

Emma Ziegellaub Eichler - edz2103@columbia.edu

November 9, 2012

Problem 1

a

```
french = alignment = []
forall 1 ≤ n ≤ m
    max = 0
    forall 0 ≤ k ≤ l
        q = q(k|n, l, m)
        forall f ∈ F
            if t(f|e_k) * q > max
                max = t(f|e_k) * q
                french[n - 1] = f
                alignment[n - 1] = k
```

Return *french*, *alignment*

The key observation to writing an efficient algorithm for this problem is that each French word f_i and alignment variable a_i are independent of all other French words f_j and alignment variables a_j , $i \neq j$, so we need only the maximum likelihood French word and alignment for each position n , $1 \leq n \leq m$, in the French sentence. There are $|F|$ possible French words and $l + 1$ possible alignments (since the English sentence is of length l , and there is also the *null* alignment). Thus, this algorithm is $O(m(l + 1)|F|)$.

Out of interest:

We in fact may improve on this slightly if we assume the number of non-zero probability translations for each English word is considerably less than $|F|$. Then by hashing the possible translations for each word in the English sentence, we can avoid iterating over all of F more than once by initializing a translations hash for the English sentence:

translations = *defaultdict(set())* // This is Python for a hash that provides a default value of an empty set $()$ for each uninitialized key

```
forall f ∈ F
    forall 1 ≤ k ≤ l
        if t(f|e_k) ≠ 0
            translations[e_k].add(f)
```

Then we modify the algorithm slightly:

```

french = alignment = []
forall 1 ≤ n ≤ m
    max = 0
    forall 0 ≤ k ≤ l
        q = q(k|n, l, m)
        forall f ∈ translations[ek]
            if t(f|ek) * q > max
                max = t(f|ek) * q
                french[n - 1] = f
                alignment[n - 1] = k

```

Return *french*, *alignment*

We now have an algorithm that runs in $O(l|F| + m(l+1)\max_{1 \leq k \leq l} |\text{translations}[e_k]|) < O(m(l+1)|F|)$ since we assume $\max_{1 \leq k \leq l} |\text{translations}[e_k]| < |F|$. \square

b

```

french = []
forall 1 ≤ n ≤ m
    max = 0
    forall f ∈ F
        temp = 0
        forall 0 ≤ k ≤ l
            temp+ = t(f|ek) * q(k|n, l, m)
        if temp > max
            max = temp
            french[n - 1] = f

```

Return *french*

The key observation to writing an efficient algorithm for this problem is that each French word f_i is independent of all other French words f_j , $i \neq j$, so we need only the maximum likelihood French word for each position n , $1 \leq n \leq m$, in the French sentence. There are $|F|$ possible French words and $l + 1$ possible alignments to sum over (since the English sentence is of length l , and there is also the *null* alignment). Thus, this algorithm is $O(m(l+1)|F|)$.

Out of interest:

We in fact may improve on this if we assume the number of non-zero probability translations for each English word is considerably less than $|F|$. Then by hashing the possible translations for each word in the English sentence, we can avoid iterating over all of F more than once by initializing a translations hash for the English sentence:
translations = *defaultdict(set())* // This is Python for a hash that provides a default value of an empty set () for each uninitialized key

```

forall f ∈ F
    forall 1 ≤ k ≤ l

```

$ift(f|e_k) \neq 0$
 $translations[f].add(k)$

Then we modify the algorithm slightly:

```
french = []
forall 1 ≤ n ≤ m
  max = 0
  forall f ∈ translations
    temp = 0
    forall k ∈ translations[f] ∩ {0}
      temp+ = t(f|e_k) * q(k|n, l, m)
    if temp > max
      max = temp
  french[n - 1] = f
```

Return *french*

We now have an algorithm that runs in $O(l|F| + m|translations| \max_{f \in translations} |translations[f] \cap \{0\}|)$. $O(m(l+1)|F|)$ since we assume $|translations| \ll |F|$ (and $\max_{f \in translations} |translations[f] \cap \{0\}| \leq l + 1 \forall f$). \square

c

By searching for $argmax_e p(f|e)PLM(e)$ instead of $argmax_e p(e|f)$, we may reuse our t and q parameters $t(f|e)$ and $q(a_j|j, l, m)$. Were we to search for $argmax_e p(e|f)$, we would have to calculate a new set of t and q parameters, $t(e|f)$ and $q(a_j|j, m, l)$. Thus, using only one set of translation parameters, we may translate in both directions, assuming we have a language model for one of the languages (which is a reasonable assumption for the native language).

2

Initialize $temp(u, 1) = q(u|0, 1, l, m) * t(f_1|e_u) \forall 0 \leq u \leq l$.

$\forall 1 \leq k \leq m$

$\forall 0 \leq v \leq l$

$temp(v, k) = \max_{0 \leq u \leq l} q(v|u, k, l, m) * t(f_k|e_v) * temp(u, k - 1)$

$bp(v, k) = \operatorname{argmax}_{0 \leq u \leq l} q(v|u, k, l, m) * t(f_k|e_v) * temp(u, k - 1)$

$a_m = \operatorname{argmax}_{0 \leq v \leq l} temp(v, m)$

$\forall 2 \leq k \leq m$

$a_{k-1} = bp(a_k, k)$

Return a_1, \dots, a_m

Given every possible alignment at every position k , $1 \leq k \leq m$ in the French sentence, we find the maximum likelihood preceding alignment (that is, the maximum likelihood alignment for position $k - 1$). We then find the maximum likelihood alignment for the last position, m , and use backpointers to recover the maximum likelihood sequence of alignments. Since we use a first-order Markov assumption (each alignment is dependent only on the alignment directly preceding it), the algorithm runs in $O(m(l + 1)^2) = O(ml^2)$ time. (The initialization and backpointer recovery run respectively in $O(l + 1) = O(l) < O(ml^2)$ and $O(l + m) < O(ml^2)$ time.) \square

3

Given input $x = x_1 \dots x_l$, initialize $temp(0, 0) = 0$

$\forall 1 \leq e_0 \leq l$

$\forall 1 \leq s_0 \leq e_0$

if $(s_0, e_0, x_{s_0} \dots x_{e_0}) \in P$

$temp(s_0, e_0) = g(s_0, e_0, x_{s_0} \dots x_{e_0}) + \max_{1 \leq s \leq s_0-1} temp(s, s_0 - 1)$

$bp(s_0, e_0) = \operatorname{argmax}_{1 \leq s \leq s_0-1} temp(s, s_0 - 1)$

$p_0 = (\operatorname{argmax}_{1 \leq s_0 \leq l} temp(s_0, l), l, x_{s_0} \dots x_l)$

$s = s_0, i = 0, e = l$

while $e > 0$ $p_{i+1} = (bp(s, e), s - 1, x_{bp(s, e)} \dots x_{s-1})$

$i ++, s = bp(s, e), e = s - 1$

Return p_{i-1}, \dots, p_0

Given every possible phrase p ending at every position e_0 , $1 \leq e_0 \leq l$ in the sentence, we find the maximum likelihood preceding phrase (that is, the maximum likelihood phrase ending at position $s(p) - 1$). We then find the maximum likelihood phrase ending in the last position, l , and use backpointers to recover the maximum likelihood sequence of phrases. The algorithm runs in $O(l^2)$ time. (The backpointer recovery is linear in the number of phrases in the maximum likelihood derivation, which is a maximum of l .) \square